

BioBridge OS

Event: QBI Hackathon 2026

Category: Business Proposal

Author: Karen Wu

1. Executive Summary

BioBridge OS is a low-code data transformation platform purpose-built for the handoff between wet-lab biologists and computational biologists. Wet-lab teams generate the raw experimental data that fuels every downstream discovery; computational teams depend on that data being clean, consistent and well-documented before any real analysis can begin. Today, that handoff process is broken and costing skilled computational staff an estimated 50–60% of their working time on cleaning rather than analysis and in turn costing wet-lab staff their autonomy, since they cannot resolve a data issue without someone else's time and code.

BioBridge OS closes this gap with a friendly, click-driven interface that lets a bench scientist resolve data inconsistencies themselves without writing a line of code or a regular expression knowledge while producing a fully reproducible Python script and an immutable, attributable audit trail for every decision made. A working functional demo has been built and validated end-to-end: raw file upload to a downloaded, reproducible cleaning script in under three minutes, with zero code written by the user.

We are seeking seed funding to move from a validated functional demo to a pilot-ready product with a small number of design-partner labs, and to build the schema/protocol layer that makes detection deterministic, auditable, and reusable across an entire lab's body of work.

2. The Problem

2.1 A structural, not a personal, failure

After fully understanding the problems both the bench scientist and computational scientists face at this handoff process point to the same root cause from two different vantage points. The computational biologist receives spreadsheets with inconsistent category naming (“Control_Rep1” vs. “ctrl_r2” vs. “CONTROL_REP1”), numeric columns contaminated with instrument-reported strings (“ERROR”, “LOW_QUALITY”, “<LOD”), and missing replicates with no flag. The wet-lab biologist, meanwhile, has no tool to see these problems before sending a file, no shared vocabulary for what counts as “wrong,” and no visibility into what was changed once her data leaves her hands.

Neither persona is at fault. The instrument software has no concept of experimental design; the spreadsheet has no schema; and no existing tool sits at the intersection of “easy enough for a bench scientist” and “rigorous enough for a computational pipeline.”

2.2 The cost, quantified

- 50–60% of computational biologist time spent on data cleaning rather than analysis, by team estimate.
- 2–4 business days of round-trip turnaround for a single data-quality clarification, based on observed handoff cycles.
- Zero existing audit trails in the majority of observed workflows and cleaning decisions are made and not recorded, creating downstream disputes about what a result actually reflects.

3. Market Opportunity

BioBridge OS sits within the broader laboratory informatics and scientific data management software market, which multiple independent industry analyses place in a sustained double-digit growth phase. Recent market research estimates the global scientific data management system software market at roughly \$3.0 billion in 2026, growing toward \$9.4 billion by 2035, while the broader laboratory informatics category like LIMS, ELN and adjacent data tools is estimated in the \$2–3 billion range for LIMS alone and tracking toward continued high-single to double-digit annual growth. The wider life sciences software market, which includes these laboratory data categories, is estimated at roughly \$19–22 billion in 2026 and forecast to approach \$40–45 billion by the early 2030s across multiple analyst estimates.

Two structural trends support BioBridge OS specifically:

1. Cloud and SaaS adoption is now the dominant deployment model across this market, with multiple analyses placing cloud-based deployment above 50% share and growing fastest — favoring a lightweight, browser-based tool over heavier on-premise informatics suites.
2. Consolidation at the top of the market (e.g., large-scale acquisitions of established platforms) is leaving a gap at the lower end, where small and mid-sized labs that cannot justify a full LIMS/ELN deployment, but still face the exact wet-lab/dry-lab handoff problem BioBridge OS is built to solve.

This is not a claim that BioBridge OS competes for the full laboratory informatics market on day one. It is evidence that the category is well-funded, growing and underserved at the specific, narrow intersection this product targets: the handoff moment itself, rather than the broader system-of-record problem incumbents already address.

4. Competitive Landscape

No incumbent we have evaluated sits directly at the data-transformation layer of the wet-lab/dry lab handoff. Each adjacent category has a structural reason it does not and will not easily address this gap:

Category	Representative players	Strength	Gap relative to BioBridge OS
Electronic Lab Notebooks	Benchling, IDBS	Excellent manual tracking and sample recording	Rigid schemas; no native code export; not built for ad hoc cleaning
Bio-Compute Orchestrators	LatchBio, DNAnexus	Powerful pipeline orchestration	Assumes data already cleaned before entry
Enterprise Data Cleansing	Databricks / Bamboolib, OpenRefine	Strong low-code transformation, similar UI philosophy	Entirely agnostic to biology; no domain vocabulary or attribution model
Traditional LIMS	Thermo Fisher SampleManager, LabWare, Dotmatics	Deep, validated, regulated systems of record	Heavy, expensive, long deployment cycles; not designed for self service cleaning by bench staff

BioBridge OS's closest functional analog for its core interaction model is OpenRefine's clustering-based cleanup — validating that the underlying technique (cluster, suggest, click-to merge) is proven. What does not exist anywhere in this landscape is that interaction model paired with a biology-aware vocabulary, an attributable audit trail, and reproducible code export, in one self-service tool.

5. The Solution: BioBridge OS

5.1 Core product principles

- Regex knowledge is not required from the bench scientist user and messy data is resolved through cluster cards, judgment-call cards, and toggles, not just pattern syntax.
- Every cleaning decision is attributed, timestamped and reasoned. This removes ambiguity and blame from the handoff relationship is the emotional and functional core of the product.
- Every visual action has a 1:1 live counterpart in generated, reproducible Python/Pandas code if a more technically advanced user wants to confirm what is happening under the hood.
- Detection is deterministic wherever a protocol/schema exists, and gracefully falls back to fuzzy, human-resolved suggestions wherever it doesn't so the system gets smarter and faster per lab over time, without ever blocking on a missing rule.

5.2 How it works, end to end

1. A wet-lab scientist uploads a raw export grid (e.g. CSV, Excel, ODS, TSV, or pasted data directly from their instrument's output).
2. BioBridge OS automatically surfaces every inconsistency: casing/category divergence, non-numeric values in numeric columns, missing replicates, missing values, and whitespace artifacts, eliminating the manual scanning usually required.
3. They resolve each issue with a click: merging a cluster of equivalent category names, or answering a plain-language judgment-call card ("what happened to this value?") for an ambiguous reading.
4. Every resolution is logged to an immutable audit trail: what changed, why, who decided it, and how long it took; this is also reflected instantly in a live, reproducible Python script.
5. A computational biologist receives a clean dataset, a script that reproduces it exactly, and a full record of every judgment call made along the way and it eliminates both the guesswork and the blame that previously defined this handoff.

5.3 What makes this defensible, not just convenient

The schema/protocol layer is the product's structural moat. Once a computational biologist defines a protocol template for a given assay (e.g. the expected columns, types, and allowed values) for every future upload against that protocol is validated deterministically, with zero noise for already-known variants. Every new resolution a bench scientist makes can be promoted back into that protocol, so the system's accuracy compounds with usage inside a given lab. This creates real switching costs and genuine data-network effects at the level of an individual lab's accumulated protocols which is a structural advantage a generic data-cleaning tool cannot replicate without the same domain-specific schema investment.

6. Product Validation

A functional, end-to-end demo has been built and tested against multiple realistic raw data scenarios, including a 96-well plate viability assay, a multi-day batch-consolidated export, a qPCR gene-expression export, and a multi-sheet Excel workbook. Demonstrated and measured results from this build:

Metric	Demonstrated result
Time from raw upload to downloaded, reproducible script	Under 3 minutes
Script-to-cleaned-preview reproducibility	100%, the exported script reproduces the in-app cleaned dataset exactly
Regex auto-generation in schema	Verified across the full resolution flow
Audit trail completeness	Every resolved anomaly produces exactly one attributed, timestamped, immutable log entry

This validation was conducted using synthetic but realistic data modeled directly on observed real-world export formats and failure modes, not idealized inputs. Note that the included demo datasets deliberately embed the same casing drift, instrument error codes and missing-replicate patterns described in Section 2.

7. Proposed Business Model

BioBridge OS is structured as SaaS, priced per seat with a usage tier for dataset volume — consistent with the dominant cloud/SaaS deployment trend already established across this market category (Section 3). Two primary customer motions are anticipated:

- **Direct-to-lab:** individual academic and biotech research labs, priced per seat per month, targeting teams currently underserved by full LIMS/ELN deployments.
- **Platform/embed:** integration into existing LIMS, ELN, or bio-compute orchestration platforms as a transformation layer, which is a natural partnership motion given the explicit gap identified in Section 4 for each of those categories.

8. Roadmap

The current functional demo validates the core “raw” data handoff workflow. The following near term roadmap items extend the product's defensibility and total addressable workflow, prioritized by the value they unlock:

8.1 Multi-batch ingestion

Accept multiple files in a single upload session and reconcile them against one schema automatically which turns the multi-day batch consolidation problem (Section 2) from a manual, error-prone reconciliation into a single guided import.

8.2 Universal instrument integration

A vendor-agnostic ingestion layer that normalizes exports from any plate reader, qPCR system, or imaging platform which removes the need for any scientist to learn a new export format every time their lab adds an instrument.

8.3 Automated continuous audit history via version control

Git-backed versioning of every dataset and protocol change, giving teams a permanent, queryable record of who changed what, when, and why; extending the existing audit trail from a per-session log into permanent institutional memory.

8.4 Already in the core build (not roadmap, included today)

- Protocol/schema authoring layer, allowing deterministic, lab-specific validation rules defined once and reused (Section 5.3).
- An explicit, opt-in “Advanced Mode” for technical users who want pattern-based rules, without compromising the default no-code experience for bench scientists.

9. Go-to-Market Strategy

1. **Design-partner pilots:** 3–5 academic or biotech labs already experiencing the wet-lab/dry lab handoff pain described in Section 2, recruited directly through existing life-sciences research networks.
2. **Land-and-expand within institutions:** a successful pilot in one lab group is a natural introduction to adjacent labs within the same core facility or department.
3. **Partnership exploration with adjacent platforms:** given the explicit gaps identified in Section 4 (ELNs, orchestrators, enterprise cleansing tools), an embedded-transformation layer partnership is a credible parallel path to direct adoption.

10. Funding Ask & Use of Funds

We are raising a seed round to fund the next 12–18 months: converting the validated functional demo into a pilot-ready product, running design-partner pilots, and building the schema/protocol and multi-batch capabilities described in Section 8. Indicative use of funds:

Category	Allocation	Notes
Engineering (schema layer, multi-batch ingestion, instrument	[TBD]%	Core roadmap items, Section 8

integration)		
Design-partner pilots & customer success	[TBD]%	3–5 lab pilots, Section 9
Go-to-market & early commercial hires	[TBD]%	Direct-to-lab and partnership motions
Operations & compliance runway	[TBD]%	Including future data handling/compliance review ahead of any regulated data deployment

Specific raise amount, valuation, and allocation percentages are intentionally presented as placeholders above. This will be finalized alongside cap table and runway modeling specific to the founding team's structure, which falls outside the scope of this product-and-market document.

11. Risks & Mitigations

Risk	Mitigation
Incumbent LIMS/ELN vendors add similar cleaning UI	Schema/protocol layer creates per-lab data-network effects that are slow for a generalist incumbent to replicate; embed-partnership motion turns potential competitors into distribution channels instead
Pilot labs require on-premise/regulated deployment before adopting a cloud tool	MVP's fully client-side architecture (no backend, no persistence) is itself a low-friction starting point; a compliance review is explicitly scoped as a prerequisite before any regulated/PHI-adjacent deployment, not an afterthought
Bench scientists resist a new tool in an already-busy workflow	Core design principle is zero training overhead: no code, no regex knowledge, plain-language judgment call cards; validated in functional demo at under 3 minutes end-to-end
Schema/protocol authoring becomes a maintenance burden for computational staff	Promote-to-protocol flow lets the schema grow from real bench usage rather than requiring upfront authoring of every possible variant

12. Closing

BioBridge OS addresses a problem that is well-documented, costly and currently unowned by any existing platform: the moment data crosses from the bench to the pipeline. The functional demo proves the core mechanic works: a non-technical user can resolve real data inconsistencies in minutes, with full attribution, producing genuinely reproducible code. The roadmap items in Section 8 extend that mechanic into a defensible, compounding asset at the level of an individual lab's accumulated knowledge.

We welcome the opportunity to walk through a live demonstration of the current build and to discuss pilot structure with prospective design-partner labs.